

Services

Search for services, features, blogs, docs, and more

[Alt+S]

New EC2 Experience

Help us shape the future

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Instances (1/2)

Search

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4	Elastic IP	IPv6 IPs	Monitoring	Security gro
Web Server 1	i-05634b2f96552f9a4	running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	3.208.24.127	-	-	disabled	Web Secur
Bastion Host	i-088feef315d9d7a0	running	t2.micro	Initializing	No alarms	us-east-1a	-	54.242.231.218	-	-	disabled	c56851a975

Instance: i-05634b2f96552f9a4 (Web Server 1)

Details

Security

Networking

Storage

Status checks

Monitoring

Tags

Instance summary

Instance ID

i-05634b2f96552f9a4 (Web Server 1)

Public IPv4 address

3.208.24.127 | open address

Private IPv4 addresses

10.0.0.14

Instance state

running

Public IPv4 DNS

-

Elastic IP addresses

-

AWS Compute Optimizer finding

User: aws-iam-sa:033401269507:assumed-role/vocalabs/user-2025863-Saitwadekar_Valay_Angar is not authorized to perform: compute-optimizer:GetEnrollmentStatus on resource: * because no identity-based policy allows the compute-optimizer:GetEnrollmentStatus action

Retry

Auto Scaling Group name

-

Instance details

Platform

Amazon Linux (Inferred)

AMI ID

ami-065efef2c739b613b

AMI name

amazon2-ami-hem-2.0.20220606.1-v86_64-gp2

Launch time

Tue Jul 05 2022 19:53:01 GMT+0530 (India Standard Time) (3 minutes)

Monitoring

disabled

Termination protection

Disabled

AMI location

amazon/amazon2-ami-hem-2.0.20220606.1-v86_64-gp2

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclate/user/2025863-Satwadekar_Valay_Angar @ 0354-0126-9507

EC2 > Instances > i-05634b2f9655279a4 > Create image

Create image

An image (also referred to as an AMI) defines the programs and settings that are applied when you launch an EC2 instance. You can create an image from the configuration of an existing instance.

Instance ID

i-05634b2f9655279a4 (Web Server 1)

Image name

WebServerAMI

Maximum 127 characters. Can't be modified after creation.

Image description - optional

Lab AMI for Web Server

Maximum 255 characters

No reboot

☒ Enable

Instance volumes

Volume type	Device	Snapshot	Size	Volume type	IOPS	Throughput	Delete on termination	Encrypted
EBS	/dev/sd...	Create new snapshot fr...	8	EBS General Purpose S...	100		<input checked="" type="checkbox"/> Enable	<input type="checkbox"/> Enable

Add volume

During the image creation process, Amazon EC2 creates a snapshot of each of the above volumes.

Tags - optional

A tag is a label that you assign to an AWS resource. Each tag consists of a key and an optional value. You can use tags to search and filter your resources or track your AWS costs.

☒ Tag image and snapshots together

Tag the image and the snapshots with the same tag.

☐ Tag image and snapshots separately

Tag the image and the snapshots with different tags.

No tags associated with the resource.

Add new tag

You can add up to 50 more tags.

Cancel

Create image

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vocalta/user/2025863-Saitowdakar_Valay_Angar @ 0354-0126-9507

New EC2 Experience

Successfully created ami-0a51114d58a821158 from instance i-05634b2f9655279a4.

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Instances (1/2)

Search

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP	IPv6 IPs	Monitoring	Security gro
Web Server 1	i-05634b2f9655279a4	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	3.208.24.127	-	-	disabled	Web Secur
Bastion Host	i-088feef315b9d97a0	Running	t2.micro	Initializing	No alarms	us-east-1a	-	54.242.231.218	-	-	disabled	c56851a975

Instance: i-05634b2f9655279a4 (Web Server 1)

Details

Security

Networking

Storage

Status checks

Monitoring

Tags

Instance summary

Instance ID

i-05634b2f9655279a4 (Web Server 1)

IPv6 address

-

Hostname type

IP name: ip-10-0-0-14.ec2.internal

Answer private resource DNS name

-

Auto-assigned IP address

3.208.24.127 (Public IP)

Public IPv4 address

3.208.24.127 | open address

Instance state

Running

Private IP DNS name (IPv4 only)

ip-10-0-0-14.ec2.internal

Instance type

t2.micro

VPC ID

vpc-0a978046149914b5d (IaB VPC)

Subnet ID

subnet-0a0b8d8df33899ab (Public Subnet 1)

Private IPv4 addresses

10.0.0.14

Public IPv4 DNS

-

Elastic IP addresses

-

AWS Compute Optimizer finding

User: aws-ec2:035401269507:assumed-role/vocalta/user/2025863-Saitowdakar_Valay_Angar is not authorized to perform: compute-optimizer:GetEnrollmentStatus on resource: * because no identity-based policy allows the compute-optimizer:GetEnrollmentStatus action

Retry

Auto Scaling Group name

-

Waiting for us-east-1-prod-pr-analytics.console.aws.a2z.com...

© 2022, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

EC2 > Target groups > Create target group

Step 1
Specify group details

Step 2
Register targets

Register targets

This is an optional step to create a target group. However, to ensure that your load balancer routes traffic to this target group you must register your targets.

Available instances (2)

Filter resources by property or value

Instance ID	Name	State	Security groups	Zone	Subnet ID
i-05634b2f96552f9a4	Web Server 1	running	Web Security Group	us-east-1a	subnet-0cddb8daf33899ab
i-088f6ef3158d9d7a0	Bastion Host	running	c56851a975485f235759911e0334012689507-BastionSecurityGroup-1E0VNE2ZYHWPV	us-east-1a	subnet-0cddb8daf33899ab

0 selected

Ports for the selected instances

Ports for routing traffic to the selected instances.

80

1-65535 (separate multiple ports with comma)

Include as pending below

Review targets

Targets (0)

All

Filter resources by property or value

Remove all pending

< 1 >

Remove	Health status	Instance ID	Name	Port	State	Security groups	Zone	Subnet ID
No instances added yet								

Specify instances above, or leave the group empty if you prefer to add targets later.

0 pending

Cancel Previous Create target group

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Successfully created target group: LabGroup

EC2 > Target groups

Target groups (1/1)

Search or filter target groups

Name	ARN	Port	Protocol	Target type	Load balancer	VPC ID
LabGroup	arn:aws:elasticloadbalancing:us-east-1:031401369507:targetgroup/LabGroup/64f5d35af36d0e	80	HTTP	Instance	None associated	vpc-0d97904d149914b5d

Target group: LabGroup

Details | Targets | Monitoring | Health checks | Attributes | Tags

Details

arn:aws:elasticloadbalancing:us-east-1:031401369507:targetgroup/LabGroup/64f5d35af36d0e

Target type	Protocol : Port	Protocol version	VPC
Instance	HTTP: 80	HTTP1	vpc-0d97904d149914b5d
IP address type	Load balancer		
IPv4	None associated		

Total targets	Healthy	Unhealthy	Unused	Initial	Draining
0	0	0	0	0	0

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vocalata/user/2025863-Satwadekar_Valay_Angar @ 0354-0126-9507

Listener HTTP:80

Remove

Protocol HTTP

Port 80

Default action info

Forward to LabGroup

Target type: Instance, IPv4

HTTP

Create target group

Add listener

Add-on services - optional

Additional AWS services can be integrated with this load balancer at launch. You can also add these and other services after your load balancer is created by reviewing the "Integrated Services" tab for the selected load balancer.

AWS Global Accelerator info

Create an accelerator to get static IP addresses and improve the performance and availability of your applications. Additional charges apply

Tags - optional

Consider adding tags to your load balancer. Tags enable you to categorize your AWS resources so you can more easily manage them. The 'Key' is required, but 'Value' is optional. For example, you can have Key = production-webserver, or Key = webserver, and Value = production.

Summary

Review and confirm your configurations. Estimate cost

Basic configuration

LabELB

Internet-facing

IPv4

Security groups

Web Security Group

sg-0830042c294826509

Network mapping

VPC vpc-0a97804d149914b5d

Lab VPC

us-east-1a

subnet-fc0d8d0aef13899e6

Public Subnet 1

us-east-1b

subnet-038220e24f48d71e5

Public Subnet 2

Listeners and routing

HTTP:80

defaults to

LabGroup

Add-on services

Name

Tags

Name

Attributes

Certain default attributes will be applied to your load balancer. You can view and edit them after creating the load balancer.

Cancel

Create load balancer

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclate/user2025863-Satwadekar_Valay_Angar @ 0354-0126-9907

Successfully created load balancer: **LabELB**

Note: It might take a few minutes for your load balancer to be fully set up and ready to route traffic. Targets will also take a few minutes to complete the registration process and pass initial health checks.

EC2 > Load balancers > Create Application Load Balancer

Create Application Load Balancer

Suggested next steps

- Review, customize, or enable attributes for your load balancer and listeners using the **Description and Listeners** tabs within **LabELB**.
- Discover other services that you can integrate with your load balancer. Visit the **Integrated services** tab within **LabELB**.

View load balancer

Feedback

Looking for language selection? Find it in the new **Unified Settings**

© 2022, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

New EC2 Experience

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Create Load Balancer

Actions

search - LabelB

1 to 1 of 1

Name	DNS name	State	VPC ID	Availability Zones	Type	Created At	Monitoring
LabelB	LabelB-1555355068-us-east-1.elb.amazonaws.com	Provisioning	vpc-0d97804d149914b5d	us-east-1b, us-east-1a	application	July 5, 2022 at 8:02:18 PM ...	

Load balancer: LabelB

Description

Listeners

Monitoring

Integrated services

Tags

Basic Configuration

NameLabelB

ARNarn:aws:elasticloadbalancing:us-east-1:933401269507:loadbalancer:app/LabelB:cs88078631c02a6

DNS nameLabelB-1555355068-us-east-1.elb.amazonaws.com

StateProvisioning

Typeapplication

Schemeinternet-facing

ID address typetpv4

VPCvpc-0d97804d149914b5d

Availability Zonessubnet-025229e24646d71d5 - us-east-1b

Hosted zoneZ:16SXDOTRQ7X7K

Creation timeJuly 5, 2022 at 8:02:18 PM UTC+5:30

Security groupsig-0830042c294828509, Web Security Group

Attributes

Deletion protectionDisabled

Idle timeout60 seconds

HTTP/2Enabled

Drain timeout modeDefensive

Feedback

Looking for language selector? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vociata/user2025863-Satwadekar_Valay_Angar @ 0354-0126-9907

Free tier eligible customers can get up to 30 GB of EBS storage. [Learn more](#) about free usage tier eligibility and usage restrictions.

Security groups

info

Assign a security group

Create a new security group

Select an existing security group

Security groups

Copy to new

View rules

Search security groups

< 1 >

Security group ID	Name	VPC ID	Description
sg-0830042c294826520	Web Security Group	vpc-0d97804d149914b5d	Enable HTTP access
sg-0916af350b99c39ec	c56851a975485c2357599f1w033401269507-RastorSecurityGroup-1EDVHE2YIMwPV	vpc-0d97804d149914b5d	Enables SSH access.
sg-00bebdba03db88f83	DB Security Group	vpc-0d97804d149914b5d	DB Security Group
sg-053e5ee9bc2354ef0	default	vpc-0d88f43fb6ba3f58e	default VPC security group
sg-048f244fd2109c2a	default	vpc-0d97804d149914b5d	default VPC security group

Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Key pair (login)

info

Key pair options

Choose an existing key pair

Existing key pair

vockey

I acknowledge that I have access to the selected private key file (vockey.pem), and that without this file, I won't be able to log into my instance.

Cancel

Create launch configuration

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

Services

Search for services, features, blogs, docs, and more

[Alt+S]

New EC2 Experience

Help us shape the future

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Successfully created launch configuration: LabConfig

EC2 > Launch configurations

Launch configurations (1/1)

Search launch configurations

Name	AMI ID	Instance type	Spot price	Creation time
LabConfig	ami-0de91f1d59821158	t2.micro	-	Tue Jul 05 2022 20:04:59 GMT+0530 (India Standard Time)

Launch configuration: LabConfig

Details

Copy launch configuration

AMI ID	ami-0de91f1d59821158	Instance type	t2.micro	IAM instance profile	-
Kernel ID	-	Key name	vockey	Monitoring	true
EBS optimized	false	Security groups	sg-0830042c294826509	Spot price	-
Create time	Tue Jul 05 2022 20:04:59 GMT+0530 (India Standard Time)	RAM disk ID	-	IP address type	Default
Metadata accessible	-	Token hop limit	-	Metadata version	-
User data	-				

Storage (volumes)

Search block devices

Block device	Size (GiB)	Type	IOPS	Throughput	Delete on termination	Encrypted
/dev/xvda	8	gp2	-	-	true	false

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclate/user/20225863-Satwadekar_Valay_Angur @ 0354-0126-9507

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Review

Choose launch template or configuration

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group. If you currently use launch configurations, you might consider migrating to launch templates.

Name

Auto Scaling group name

Enter a name to identify the group.

Lab Auto Scaling Group

Must be unique to this account in the current Region and no more than 255 characters.

Launch configuration

Switch to launch template

Launch configuration

Choose a launch configuration that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.

LabConfig

Create a launch configuration

Launch configuration

LabConfig

AMI ID

ami-0de91fd59a821158

Date created

Tue Jul 05 2022 20:04:58 GMT+0530 (India Standard Time)

Security groups

sg-0830042c294826509

Instance type

t2.micro

Key pair name

vockey

Cancel

Next

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vocalata/user/2025863-Satwadekar_Valay_Angar @ 0354-0126-9507

EC2 > Auto Scaling groups > Create Auto Scaling group

Step 1
Choose launch template or configuration

Step 2
Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Review

Choose instance launch options

Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.

Network

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-0d57804d149914b5d (Lab VPC)

Create a VPC

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

us-east-1a | subnet-076f5f874586586e (Private Subnet 1)

us-east-1b | subnet-07b99acaf1e14818 (Private Subnet 2)

Create a subnet

Cancel

Previous

Skip to review

Next

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vocalata/user/20250631-Satwadekar_Valay_Angur @ 0354-0126-9507

Choose instance launch options

Step 3 (optional)
Configure advanced options

Step 4 (optional)
Configure group size and scaling policies

Step 5 (optional)
Add notifications

Step 6 (optional)
Add tags

Step 7
Preview

Load balancing - optional

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

No load balancer
Traffic to your Auto Scaling group will not be forwarded by a load balancer.

Attach to an existing load balancer
Choose from your existing load balancers.

Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

Attach to an existing load balancer

Select the load balancers that you want to attach to your Auto Scaling group.

Choose from your load balancer target groups
This option allows you to attach Application, Network, or Gateway Load Balancers.

Choose from Classic Load Balancers

Existing load balancer target groups

Only Amazon target groups that belong to the same VPC as your Auto Scaling group are available for selection.

Select target groups

LabGroup | HTTP
Application Load Balancer: LabELB

Health checks - optional

Health check type
EC2 Auto Scaling automatically replaces instances that fail health checks. If you enabled load balancing, you can enable ELB health checks in addition to the EC2 health checks that are always enabled.

EC2

ELB

Health check grace period

The amount of time until EC2 Auto Scaling performs the first health check on new instances after they are put into service.

300 seconds

Additional settings - optional

Monitoring
Enable group metrics collection within CloudWatch

Default instance warmup

The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.

Enable default instance warmup

Cancel

Previous

Skip to review

Next

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2025, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

Step 2

Choose instance launch options

Step 3 (optional)

Configure advanced options

Step 4 (optional)

Configure group size and scaling policies

Step 5 (optional)

Add notifications

Step 6 (optional)

Add tags

Step 7

Review

Group size - optional

Specify the size of the Auto Scaling group by changing the desired capacity. You can also specify minimum and maximum capacity limits. Your desired capacity must be within the limit range.

Desired capacity

2

Minimum capacity

2

Maximum capacity

6

Scaling policies - optional

Choose whether to use a scaling policy to dynamically resize your Auto Scaling group to meet changes in demand.

Target tracking scaling policy

Choose a desired outcome and have it to the scaling policy to add and remove capacity as needed to achieve that outcome.

None

Scaling policy name

LabScalingPolicy

Metric type

Average CPU utilization

Target value

60

Instances need

300

seconds warm up before including in metric

Disable scale in to create only a scale-out policy

Instance scale-in protection - optional

Instance scale-in protection

If protect from scale in is enabled, newly launched instances will be protected from scale in by default.

Enable instance scale-in protection

Cancel

Previous

Skip to review

Next

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

voclate/user/2025863-Satwadekar_Valay_Angar @ 0354-0126-9507

Enabled

Disabled

Step 4: Configure group size and scaling policies

Edit

Group size

Desired capacity	Minimum capacity	Maximum capacity
2	2	6

Scaling policy

Target tracking scaling	Scaling policy name	Execute policy when
Policy type	LabScalingPolicy	As required to maintain Average CPU utilization at 60
Target tracking scaling		
Take the action	Instances need	Scale in
Add or remove capacity units as required	300 seconds to warm up before including in metric	Enabled

Instance scale-in protection

Instance scale-in protection

☒ Enable instance protection from scale in

Step 5: Add notifications

Edit

Notifications

No notifications

Step 6: Add tags

Edit

Tags (1)

Key	Value	Tag new instances
Name	Lab Instance	Yes

Cancel>Create Auto Scaling group

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vocalta/user/20230631-Satwadekar_Valay_Angar @ 0354-0126-9907

New EC2 Experience

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMIs

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

The old Auto Scaling groups console is no longer available. We will keep improving the new console based on your feedback.

Predictive scaling policy now supports custom metrics, which also allows you to retain metrics across Blue/Green deployments.

Lab Auto Scaling Group, 1 Scaling policy created successfully. Group metrics collection is enabled.

EC2 > Auto Scaling groups

Auto Scaling groups (1/1)

Search your Auto Scaling groups

NameLaunch template/configurationInstancesStatusDesired capacityM...M...Availability Zones

Lab Auto Scaling GroupLabConfig0Updating capacity226us-east-1a, us-east-1b

DetailsActivityAutomatic scalingInstance managementMonitoringInstance refresh

Group details

Desired capacity2Auto Scaling group nameLab Auto Scaling Group

Minimum capacity2Data createdTue Jul 05 2022 20:09:12 GMT+05:30 (India Standard Time)

Maximum capacity6Amazon Resource Name (ARN)arn:aws:autoscaling:us-east-1:1033401269507:autoScalingGroup:19166fae-e09b-49d0-a671-9c51bd52f069:autoScalingGroupName/Lab Auto Scaling Group

Launch configuration

Launch configurationLabConfigAMI IDami-0de91f1d59a821158

Instance typet2.microKey pair namevockey

Storage (volumes)/dev/xvdaSecurity groupssg-0830042c294826509

View details in the launch configuration console

Create timeTue Jul 05 2022 20:04:59 GMT+05:30 (India Standard Time)

Network

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. PrivacyTermsCookie preferences

Services

Search for services, features, blogs, docs, and more

[Alt+Q]

New EC2 Experience

Find or create your things

Instances (1/4)

info

Connect

Instance state

Actions

Launch instances

Services

Search for services, features, blogs, docs, and more

[Alt+Q]

New EC2 Experience

Find or create your things

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMI

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

EC2 > Target groups

Target groups (1/1)

info

Search or filter target groups

LabGroup

arn:aws:elasticloadbalancing:us-east-1:1093180602488026000:targetgroup/TargetGroup/1093180602488026000

80

HTTP

Instance

LabELB

vpc-0d97804d149b14b5d

Target group: LabGroup

×

Details

Targets

Monitoring

Health checks

Attributes

Tags

Registered targets (2)

Refresh

Deregister

Register targets

Filter resources by property or value

Instance ID

Name

Port

Zone

Health status

Health status details

i-0b705294ab574720e

Lab Instance

80

us-east-1a

healthy

i-09318060248802600

Lab Instance

80

us-east-1b

healthy

Feedback

Looking for language selection? Find it in the new Unified Settings.

© 2022, Amazon Web Services, Inc. or its affiliates.

Privacy

Terms

Cookie preferences



Load Test

RDS

Meta-Data	Value
InstanceId	i-0b705294ab57d72ce
Availability Zone	us-east-1a

Current CPU Load: 0%

CloudWatch

Alarms

Alarms

Alarms (1/2)

Hide Auto Scaling alarms

Clear selection

Create composite alarm

Actions

Create alarm

Search

Any state

Any type

Name	State	Last state update	Conditions	Actions
TargetTracking-Lab Auto Scaling Group-AlarmHigh-8b37c8a0-57fb-41dc-8f5b-3ebfa855ba758	OK	2022-07-05 20:12:03	CPUUtilization > 60 for 3 datapoints within 3 minutes	Actions enabled
TargetTracking-Lab Auto Scaling Group-AlarmLow-a7333c3f-6ee9-4c06-ba35-ea4f783fc79b	Insufficient data	2022-07-05 20:09:13	CPUUtilization < 54 for 15 datapoints within 15 minutes	Actions enabled

Feedback

Looking for language selection? Find it in the new [Unified Settings](#)

© 2022, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

CloudWatch

Alarms

Alarms (1/2)

Hide Auto Scaling alarms

Clear selection

Create composite alarm

Actions

Create alarm

Any state

Any type

Name	State	Last state update	Conditions	Actions
TargetTracking-Lab Auto Scaling Group AlarmLow-a7335c2f-6ee9-4c06-ba35-ea47659c799b	OK	2022-07-05 20:14:55	CPUUtilization < 54 for 15 datapoints within 15 minutes	Actions enabled
TargetTracking-Lab Auto Scaling Group AlarmHigh-8837c8a0-57bb-41dc-8f56-3e6b3358a758	OK	2022-07-05 20:12:03	CPUUtilization > 60 for 3 datapoints within 3 minutes	Actions enabled

Feedback

Looking for language selection? Find it in the new [Unified Settings](#)

© 2022, Amazon Web Services, Inc. or its affiliates. [Privacy](#) [Terms](#) [Cookie preferences](#)

[illegible]

Services

Search for services, features, blogs, docs, and more

[Alt+S]

N. Virginia

vocalta/user/2025863-Satwadekar_Valay_Angar @ 0354-0126-9507

New EC2 Experience

Web server instances

EC2 Dashboard

EC2 Global View

Events

Tags

Limits

Instances

Instance Types

Launch Templates

Spot Requests

Savings Plans

Reserved Instances

Dedicated Hosts

Scheduled Instances

Capacity Reservations

Images

AMI Catalog

Elastic Block Store

Volumes

Snapshots

Lifecycle Manager

Network & Security

Security Groups

Elastic IPs

Placement Groups

Key Pairs

Network Interfaces

Load Balancing

Load Balancers

Target Groups

Auto Scaling

Launch Configurations

Auto Scaling Groups

Successfully terminated i-05634b2f9655279a4

Instances (1/4)

Search

Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4 ...	Elastic IP	IPv6 IPs	Monitoring	Security gro
Web Server 1	i-05634b2f9655279a4	Shutting-down	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	3.208.24.127	-	-	disabled	Web Secur
Bastion Host	i-088feef315b9d97a0	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	54.242.231.218	-	-	disabled	c56851a975
Lab Instance	i-06705294a657d720a	Running	t2.micro	2/2 checks passed	No alarms	us-east-1a	-	-	-	-	enabled	Web Secur
Lab Instance	i-083b80ac02488024d0	Running	t2.micro	2/2 checks passed	No alarms	us-east-1b	-	-	-	-	enabled	Web Secur

Instance: i-05634b2f9655279a4 (Web Server 1)

Details

Security

Networking

Storage

Status checks

Monitoring

Tags

Instance summary

Instance ID

i-05634b2f9655279a4 (Web Server 1)

IPv6 address

-

Hostname type

IP name: ip-10-0-0-14.ec2.internal

Answer private resource DNS name

-

Auto-assigned IP address

3.208.24.127 (Public IP)

Public IPv4 address

3.208.24.127 | open address

Instance state

Shutting-down

Private IP DNS name (IPv4 only)

ip-10-0-0-14.ec2.internal

Instance type

t2.micro

VPC ID

vpc-0a978046149914b5d (Lab VPC)

Subnet ID

subnet-0a0b8d4ef33899ab (Public Subnet 1)

Private IPv4 addresses

10.0.0.14

Public IPv4 DNS

-

Elastic IP addresses

-

AWS Compute Optimizer finding

User: aws-ec2:035401269507:assumed-role/vocalta/user/2025863-Satwadekar_Valay_Angar is not authorized to perform: compute-optimizer:GetEnrollmentStatus on resource: * because no identity-based policy allows the compute-optimizer:GetEnrollmentStatus action

Retry

Auto Scaling Group name

-

Feedback

Looking for language selection? Find it in the new Unified Settings

© 2022, Amazon Web Services, Inc. or its affiliates. Privacy Terms Cookie preferences